

CSCI 104

Bloom Filters & Tries


Mark Redekopp

David Kempe

Sandra Batista

Set Review

- Recall the operations a set performs...
 - Insert(key)
 - Remove(key)
 - Contains(key) : bool (a.k.a. find())
- We can implement a set using
 - List
 - $O(n)$ for some of the three operations
 - (Balanced) Binary Search Tree
 - $O(\log n)$ insert/remove/contains
 - Hash table
 - $O(1)$ insert/remove/contains



Bloom Filter



- A Bloom filter is a set such that "contains()" will *quickly* answer...
 - "No" correctly (i.e. if the key is not present)
 - "Yes" with a chance of being incorrect (i.e. the key may not be present but it might still say "yes")

Bloom Filter Motivation

- Why would we want this?
 - A Bloom filter usually sits in front of an actual set/map
 - Suppose that set/map is EXPENSIVE to access
 - if set/map doesn't sit on a disk drive or another server
 - Disk/Network access = ~milliseconds
 - Memory access = ~nanoseconds
 - The Bloom filter is small enough to reside in memory for quick access and can answer quickly if the set/map on disk contains a key:
 - If it answers "No" do not search the EXPENSIVE set
 - If it answers "Yes" search the EXPENSIVE set


Bloom Filters

- A Bloom filter is...
 - A hash table of individual bits (Booleans: T/F)
 - A set of hash functions, $\{h_1(k), h_2(k), \dots, h_s(k)\}$
- Insert()
 - Apply each $h_i(k)$ to the key
 - Set $a[h_i(k)] = \text{True}$



Bloom Filters

- A Bloom filter is...
 - A hash table of individual bits (Booleans: T/F)
 - A set of hash functions, $\{h_1(k), h_2(k), \dots, h_s(k)\}$
- Contains()
 - Apply each $h_i(k)$ to the key
 - Return True if **all** $a[h_i(k)] = \text{True}$
 - Return False otherwise
 - In other words, answer is "Maybe" or "No"
 - May produce "false positives"
 - May NOT produce "false negatives"
- We will ignore removal for now



Practice

- Trace a Bloom Filter on the following operations:

- insert(0), insert(1), insert(2), insert(8), contains(2), contains(3), contains(4), contains(9)


	0	1	2	3	4	5	6	7	8	9
a	0	0	0	0	0	0	0	0	0	0

- The hash functions are
 - $h1(k) = (7k+4)\%10$
 - $h2(k) = (2k+1)\%10$
 - $h3(k) = (5k+3)\%10$
 - The table size is 10 ($m=10$).

	H1(k)	H2(k)	H3(k)	Hit?
Insert(0)	4	1	3	N/A
Insert(1)	1	3	8	N/A
Insert(2)	8	5	3	N/A
Insert(8)	0	7	3	N/A
Contains(2)	8	5	3	Yes
Contains(3)	5	7	8	Yes
Contains(4)	2	9	3	No
Contains(9)	7	9	8	No

Tries


- Goal: Achieve linear search in length of keys independent of number of keys
- Rather than each node storing a full key value, each node represents a prefix of the key
- Highlighted nodes indicate terminal locations
 - For a map we could store the associated value of the key at that terminal location
- Notice we "share" paths for keys that have a common prefix



Trie for the keys: "HE", "HEAP",
"HEAR", "HELP", "ILL",
"IN"

Tries

- To search for a key, start at the root consuming one unit (bit, char, etc.) of the key at a time
 - If highlighted node, SUCCESS
 - else FAILURE
- Examples:
 - Search for "He"
 - Search for "Help"
 - Search for "Head"
- Search takes $O(m)$ where m = length of key



A "value" type
could be stored

Application: IP Lookups

- Network routers form the backbone of the Internet
- Incoming packets contain a destination IP address (128.125.73.60)
- Routers contain a "routing table" mapping some prefix of destination IP address to output port
 - 128.125.x.x => Output port C
 - 128.209.32.x => Output port B
 - 128.209.44.x => Output port D
 - 132.x.x.x => Output port A
- Keys = Match the longest prefix
 - Keys are unique
- Value = Output port



Octet 1	Octet 2	Octet 3	Port
10000000	01111101		C
10000000	11010001	00100000	B
10000000	11010001	00101100	D
10000100			A

11

- 

Octet 1	Octet 2	Octet 3	Port
10000000	01111101		C
10000000	11010001	00100000	B
10000000	11010001	00101100	D
10000100			A



Trie Creation Algorithm

➤ Given a set of strings, S

Let Alpha be set of union of all letters in strings in S .

– The root corresponds to empty string at level 0.

For (each node at level i) /*prefix length at node is i */

For (each letter, w , in Alpha)

For (each string, s , in S)

If ($x*w$ is a prefix of s)

Add new node and edge labeled w .

New node prefix is $x*w$ of length $i+1$.

if ($x*w$ equals s) then highlight new node and add value.

If all s in S are highlighted, stop.

Trie Creation

➤ Let's construct a trie to store the set of words:

- Heap
- Helm
- Hear
- Help
- He
- In
- Ink
- Irk
- She

Trie Complexity

- What is the runtime of insert?
- What is runtime of search?
- What is size of trie?

Compressed Trie

- In our construction many nodes were redundant: only a single child along path to node
- Compressed trie is when paths of single node are combined into a node
- In compressed trie all nodes have at least one of the following properties:
 - Root node
 - Word node
 - Has at least 2 children

Compressed Trie Practice

- Let's construct a compressed trie to store the following set of words:
- Ten
 - Tent
 - Then
 - Tense
 - Tens
 - Tenth